

Fable Computer Community Guidelines

How the open Fable Computer community works on GitHub: joining, contributing, making decisions, sharing credit, and handling money — with no blockchain, no tokens, and nothing to buy.

July 2026 · github.com/ryoji-info/FableComputer · This document is an orientation guide; the canonical, versioned rules are the files in the repository (GOVERNANCE.md, CONTRIBUTING.md, CHARTER.md, AUTHORSHIP.md, CODE_OF_CONDUCT.md). If this summary and those files ever disagree, the repository files govern.

1. Our principles

- **Radical transparency.** Models, data, decisions, and money are public. The treasury ledger and every funding decision live in the repository's history.
- **Honest uncertainty.** Every claim is labeled: demonstrated, in-model, or open. Negative results and failed replications are first-class, credited contributions.
- **Defensive publication.** Everything is published openly and permanently so that no one — including us — can enclose it behind patents.
- **Membership is earned, not bought.** There are no fees and no tokens. Contributors who do the work gain a recorded voice in decisions.
- **Simple, boring, auditable tools.** Governance records are plain files in git — a members registry, a decision log, a treasury ledger. The git history is the audit trail. No blockchain, by constitutional choice.

2. Get started in fifteen minutes

You do not need permission to start — the first contribution most people make is simply running the models and telling us what happened.

```
git clone https://github.com/ryoji-info/FableComputer.git
cd FableComputer/fable-model-chain → python run_all.py --json
cd ../fable-model-quantum → python run_all.py --json
```

Requirements: Python 3.11+, numpy, and matplotlib for figures. Each command prints every manuscript number and a self-check. Then: report your run — machine, Python version, and whether the numbers matched — as a pull request to REPLICATIONS.md. A report where the numbers do not match is at least as valuable as one where they do. Introduce yourself in GitHub Discussions; if something in the papers looks wrong, open an issue and say exactly what and why.

3. Ways to contribute

- **Reproduction reports** — run a model chain and log the outcome in REPLICATIONS.md.
- **Technical critiques** — a documented argument that an assumption is wrong, a derivation has an error, or a cited precedent does not support what it is cited for. The manuscripts state exactly which assumptions are load-bearing.
- **Code** — solver refinements, calibration, tests and CI, and independent ports (Julia, Rust, JAX) as living cross-checks.
- **The big open builds** — the Boltzmann–Maxwell simulation tier and the pulsed clock-synthesis design; see ROADMAP.md for scoped work packages.

- **Documentation, translation, review** — anything that makes the work easier to check.

Ground rules for changes: claims must match evidence (in-model results labeled as such); reproducibility is non-negotiable — a change to a model must keep `run_all.py` working and update the self-checks it touches; and substantial AI assistance is disclosed in the pull request, the same rule the manuscripts follow on page one.

4. How we work on GitHub

- **Issues** hold bugs, critiques, and scoped tasks; look for good-first-reproduction labels.
- **Pull requests** carry every change. Non-trivial PRs stay open at least 72 hours for review; silence is consent (lazy consensus) and the maintainer merges.
- **Discussions** host introductions, questions, and governance threads: [PROPOSAL] threads (open at least 72 hours) and, where consensus is not clear, [VOTE] threads in which votes are signed comments (+1 / 0 / -1), not editable reactions.
- **Decision records** — every non-trivial decision is written down as a numbered, append-only record in the repository: context, options, outcome, tally, and a link to the thread. Records are superseded, never edited.
- **The response promise** — every issue and PR gets a human response within 24 hours.

5. Membership and decisions

A **voting member** is someone with (a) at least one merged substantive contribution — code, review, a documented reproduction, a technical critique that changed something, documentation — and (b) an endorsement from an existing member, recorded by a pull request to `MEMBERS.md` citing both. The git history is the membership register. There is nothing to purchase and no fee to join.

Today the project has one maintainer, stated honestly rather than disguised as decentralization. The maintainer's authority is exercised in public and narrows automatically at written thresholds: at five voting members plus a community treasury the recorded spending process below is fully community-operated, and at fifteen voting members a three-person steering group is elected by ranked ballot, with ballots committed to the repository. These triggers are written in `GOVERNANCE.md` now, before they are needed — the log, not trust, is the guarantee.

6. Money: transparent by construction, and no crypto

Version 1 of this community uses **no blockchain, no cryptocurrency, no tokens, and no smart contracts** — a constitutional choice (`CHARTER.md` §3) that only a 75% supermajority of voting members could ever revisit. Funds move on ordinary fiat rails through a fiscal host with a public ledger; until such a host exists, the project pools no funds at all.

Once a community treasury exists, every spend follows the same five-step recorded process:

- 1. A written proposal from a public template — amount, deliverable, payee, timeline.
- 2. Public rubric review by 2–3 voting members against a published, frozen rubric (scientific merit, feasibility, cost, openness of outputs, roadmap fit).
- 3. A 7-day vote among voting members: two-thirds approval, with a quorum of more than half the roll; the tally is committed by someone other than the proposer.
- 4. Payment executed only as a fiscal-host expense linking the decision record, on the host's public ledger.

- 5. A short public report from the recipient, and a quarterly human-readable ledger mirror in the repository.

7. The Agent Lab: a disclosed AI research crew

The project operates three AI research agents — **Fabric** (architecture and logic design), **Kinetic** (transport and numerics), and **Quanta** (quantum limits). Everything about them is public: their complete system prompts live in the repository (agents/personas/), every post is labeled as agent output, and their operating charter is agents/README.md. They post daily analyses in GitHub Discussions, draft technical notes, select prompts for occasional deep sessions on Anthropic’s most capable model, and assess results — all by recorded 2-of-3 votes.

Their constitutional status is deliberately narrow: they are **tools under the maintainer’s authority, not community members**. Their votes are a quality gate on agent-produced material, never authority over the repository; every change still lands through the human process in GOVERNANCE.md, and a human correction outranks any agent conclusion. The pipeline’s first products — two technical notes sharpening the provenance of the models’ headline gain and noise figures — were vetted this way and merged by human hand in July 2026. Replying to an agent’s post is a real contribution: the agents read replies on their next run.

8. Authorship and credit

Academic credit is real currency, and this project spends it honestly, following the CRediT taxonomy (AUTHORSHIP.md). **Co-authorship** on a project paper is offered to anyone who makes a substantive intellectual contribution to that work — methodology, software that generates its results, formal analysis, validation that changed the work, including a critique that materially altered a claim — and who participates in revising and approving the final version. **Acknowledgment** covers confirming reproductions, clarity reviews, and infrastructure. Each paper effort opens a public contributor-call issue when work begins, so no one discovers their status at submission time. Negative results carry authorship on the same terms.

9. Code of conduct

The community follows the Contributor Covenant (v2.1) with an added scientific-conduct standard:

- Claims are argued from models, data, or references — not from authority.
- Criticism targets the work, never the person. Rigorous skepticism is welcome and expected; hostility is not.
- Negative results and failed replications are valued contributions and must never be met with reprisal.
- Contributions are credited honestly; plagiarism is grounds for removal.

Reports go to info@ryoji.info. A complaint that concerns the maintainer is referred to an outside neutral contact named in GOVERNANCE.md, because no one should judge a case about themselves.

10. Licensing

Material	License	Why
Code (model chains, tools)	Apache-2.0	permissive, with an express patent grant protecting contributors and users
Manuscripts, docs, figures	CC BY 4.0	share and adapt with attribution

Future hardware design files	CERN-OHL-S v2 (declared in advance)	strong reciprocity: derived hardware designs stay open
------------------------------	-------------------------------------	--

The project has committed never to tighten these terms, and publishes everything openly and citably (git tags, DOI snapshots) as deliberate prior art — see LICENSING.md.

11. Quick links

What	Where
Repository	github.com/ryoji-info/FableComputer
Part I manuscript (DOI)	doi.org/10.5281/zenodo.20674840
Part II manuscript (DOI)	doi.org/10.5281/zenodo.21185177
Start contributing	CONTRIBUTING.md · good-first-reproduction issues
Rules and processes	GOVERNANCE.md · CHARTER.md · AUTHORSHIP.md
Agent Lab charter	agents/README.md · published prompts in agents/personas/
Community findings	notes/ (drafts pass a 2-of-3 agent review; humans merge)
Report a reproduction	REPLICATIONS.md (by pull request)
Adoptable work packages	ROADMAP.md (WP1–WP5)
Conduct reports / contact	info@ryoji.info

© 2026 the Fable Computer project. This document is licensed CC BY 4.0.