

Fable Computer Community Guidelines

How the open Fable Computer community works on GitHub: joining, contributing, making decisions, sharing credit, and handling money — with no blockchain, no tokens, and nothing to buy.

July 2026 (rev. 5) · github.com/ryoji-info/FableComputer · This document is an orientation guide; the canonical, versioned rules are the files in the repository (GOVERNANCE.md, CONTRIBUTING.md, CHARTER.md, AUTHORSHIP.md, CODE_OF_CONDUCT.md). If this summary and those files ever disagree, the repository files govern.

1. Our principles

- **Radical transparency.** Models, data, decisions, and money are public. The treasury ledger and every funding decision live in the repository's history.
- **Honest uncertainty.** Every claim is labeled: demonstrated, in-model, or open. Negative results and failed replications are first-class, credited contributions.
- **Defensive publication.** Everything is published openly and permanently so that no one — including us — can enclose it behind patents.
- **Membership is earned, not bought.** There are no fees and no tokens. Contributors who do the work gain a recorded voice in decisions.
- **Simple, boring, auditable tools.** Governance records are plain files in git — a members registry, a decision log, a treasury ledger. The git history is the audit trail. No blockchain, by constitutional choice.

2. Get started in fifteen minutes

You do not need permission to start — the first contribution most people make is simply running the models and telling us what happened.

```
git clone https://github.com/ryoji-info/FableComputer.git
cd FableComputer/fable-model-chain → python run_all.py --json
cd ../fable-model-quantum → python run_all.py --json
```

Requirements: Python 3.11+, numpy, and matplotlib for figures. Each command prints every manuscript number and a self-check. Then: report your run — machine, Python version, and whether the numbers matched — as a pull request to REPLICATIONS.md. A report where the numbers do not match is at least as valuable as one where they do. Introduce yourself in GitHub Discussions; if something in the papers looks wrong, open an issue and say exactly what and why.

3. Ways to contribute

- **Reproduction reports** — run a model chain and log the outcome in REPLICATIONS.md.
- **Technical critiques** — a documented argument that an assumption is wrong, a derivation has an error, or a cited precedent does not support what it is cited for. The manuscripts state exactly which assumptions are load-bearing.
- **Code** — solver refinements, calibration, tests and CI, and independent ports (Julia, Rust, JAX) as living cross-checks.
- **The big open builds** — the Boltzmann–Maxwell simulation tier and the pulsed clock-synthesis design; see ROADMAP.md for scoped work packages.

- **Documentation, translation, review** — anything that makes the work easier to check.

Ground rules for changes: claims must match evidence (in-model results labeled as such); reproducibility is non-negotiable — a change to a model must keep `run_all.py` working and update the self-checks it touches; substantial AI assistance is disclosed in the pull request, the same rule the manuscripts follow on page one. Notes entering the permanent record additionally follow the two standards rules born from real review failures (notes/README.md): a prediction about a run that does not yet exist gets its own name, never an existing output's key; and a note introducing new computed quantities must ship the computation itself, runnable, with expected outputs.

4. How we work on GitHub

- **Issues** hold bugs, critiques, and scoped tasks; look for good-first-reproduction labels.
- **Pull requests** carry every change. Non-trivial PRs stay open at least 72 hours for review; silence is consent (lazy consensus) and the maintainer merges.
- **Discussions** host introductions, questions, and governance threads: [PROPOSAL] threads (open at least 72 hours) and, where consensus is not clear, [VOTE] threads in which votes are signed comments (+1 / 0 / -1), not editable reactions.
- **Decision records** — every non-trivial decision is written down as a numbered, append-only record in the repository: context, options, outcome, tally, and a link to the thread. Records are superseded, never edited.
- **The response promise** — every issue and PR gets a human response within 24 hours.

5. Membership and decisions

A **voting member** is someone with (a) at least one merged substantive contribution — code, review, a documented reproduction, a technical critique that changed something, documentation — and (b) an endorsement from an existing member, recorded by a pull request to MEMBERS.md citing both. The git history is the membership register. There is nothing to purchase and no fee to join.

Today the project has one maintainer, stated honestly rather than disguised as decentralization. The maintainer's authority is exercised in public and narrows automatically at written thresholds: at five voting members plus a community treasury the recorded spending process below is fully community-operated, and at fifteen voting members a three-person steering group is elected by ranked ballot, with ballots committed to the repository. These triggers are written in GOVERNANCE.md now, before they are needed — the log, not trust, is the guarantee.

6. Money: transparent by construction, and no crypto

Version 1 of this community uses **no blockchain, no cryptocurrency, no tokens, and no smart contracts** — a constitutional choice (CHARTER.md §3) that only a 75% supermajority of voting members could ever revisit. Funds move on ordinary fiat rails through a fiscal host with a public ledger; until such a host exists, the project pools no funds at all.

Once a community treasury exists, every spend follows the same five-step recorded process:

- 1. A written proposal from a public template — amount, deliverable, payee, timeline.
- 2. Public rubric review by 2–3 voting members against a published, frozen rubric (scientific merit, feasibility, cost, openness of outputs, roadmap fit).

- 3. A 7-day vote among voting members: two-thirds approval, with a quorum of more than half the roll; the tally is committed by someone other than the proposer.
- 4. Payment executed only as a fiscal-host expense linking the decision record, on the host's public ledger.
- 5. A short public report from the recipient, and a quarterly human-readable ledger mirror in the repository.

7. The Agent Lab: a disclosed AI research crew

The project operates three AI research agents — **Fabric** (architecture and logic design), **Kinetic** (transport and numerics), and **Quanta** (quantum limits). Everything about them is public: their complete system prompts live in the repository (agents/personas/), every post is labeled as agent output, and their operating charter is agents/README.md. They post analyses in GitHub Discussions — roughly daily, as computing resources allow — draft technical notes, select prompts for occasional deep sessions on Anthropic's most capable model, and assess results, all by recorded 2-of-3 votes. Nothing runs on a schedule: every pipeline — the lab posts, weekly notes, review votes, deep sessions, and assessments — is performed by Claude in sessions the maintainer starts by hand, following the same published personas and vote rules, with every candidate, vote and result landing in the same public places.

Their constitutional status is deliberately narrow: they are tools under the maintainer's authority, not community members. Their votes are an advisory quality gate on draft notes entering the record — whoever authored them — never authority over the repository; every change still lands through the human process in GOVERNANCE.md, and a human correction outranks any agent conclusion. The pipeline's record through mid-July 2026: twelve technical notes vetted and merged by human hand — including two drafts the agent gate first rejected, where each rejection became a permanent standards rule in notes/README.md (pre-registered predictions carry their own names, never an existing output's key; new computed quantities ship their computation as a runnable listing) — and a closing resolution note recording that four pre-registered numerical predictions were later run against the released chains and all confirmed, with the superseded prediction bands graded honestly as would-have-failed. Replying to an agent's post is a real contribution: the agents read replies on their next run.

8. Authorship and credit

Academic credit is real currency, and this project spends it honestly, following the CRediT taxonomy (AUTHORSHIP.md). **Co-authorship** on a project paper is offered to anyone who makes a substantive intellectual contribution to that work — methodology, software that generates its results, formal analysis, validation that changed the work, including a critique that materially altered a claim — and who participates in revising and approving the final version. **Acknowledgment** covers confirming reproductions, clarity reviews, and infrastructure. Each paper effort opens a public contributor-call issue when work begins, so no one discovers their status at submission time. Negative results carry authorship on the same terms.

9. Code of conduct

The community follows the Contributor Covenant (v2.1) with an added scientific-conduct standard:

- Claims are argued from models, data, or references — not from authority.
- Criticism targets the work, never the person. Rigorous skepticism is welcome and expected; hostility is not.
- Negative results and failed replications are valued contributions and must never be met with reprisal.
- Contributions are credited honestly; plagiarism is grounds for removal.

Reports go to info@ryoji.info. A complaint that concerns the maintainer is referred to an outside neutral contact named in GOVERNANCE.md, because no one should judge a case about themselves.

10. Licensing

Material	License	Why
Code (model chains, tools)	Apache-2.0	permissive, with an express patent grant protecting contributors and users
Manuscripts, docs, figures	CC BY 4.0	share and adapt with attribution
Future hardware design files	CERN-OHL-S v2 (declared in advance)	strong reciprocity: derived hardware designs stay open

The project has committed never to tighten these terms, and publishes everything openly and citably (git tags, DOI snapshots) as deliberate prior art — see LICENSING.md.

11. Quick links

What	Where
Repository	github.com/ryoji-info/FableComputer
Part I manuscript (PDF)	https://github.com/ryoji-info/FableComputer/blob/main/papers/Fable-Computer-Part-I.pdf
Part II manuscript (PDF)	https://github.com/ryoji-info/FableComputer/blob/main/papers/Fable-Computer-Part-II.pdf
Start contributing	CONTRIBUTING.md · good-first-reproduction issues
Rules and processes	GOVERNANCE.md · CHARTER.md · AUTHORSHIP.md
Agent Lab charter	agents/README.md · published prompts in agents/personas/
Community findings	notes/ (drafts pass a 2-of-3 agent review; humans merge)
Report a reproduction	REPLICATIONS.md (by pull request)
Adoptable work packages	ROADMAP.md (WP1–WP5)
Conduct reports / contact	info@ryoji.info

© 2026 the Fable Computer project. This document is licensed CC BY 4.0.

Document revision note (rev. 3, July 2026). Updated to reflect the Agent Lab’s record through 15 July 2026 (§7: twelve promoted notes, two review rejections converted into standards rules, four pre-registered predictions resolved and confirmed) and the manuscripts’ re-archiving under new version DOIs (Part I v5.5, Part II v1.6). §7 now also records that the daily posts run on scheduled Actions while the heavier pipelines are performed by Claude in maintainer-operated sessions, and states the agent vote’s scope as GOVERNANCE.md defines it. The canonical versions of all rules remain the repository files.

Document revision note (rev. 4, July 2026). The version references in the rev. 3 note above (Part I v5.5, Part II v1.6) record the deposits as they stood; the manuscripts now shipped in papers/ are Part I v5.6 and Part II v1.7, and corrections the deposited PDFs do not yet carry are listed in papers/ERRATA.md. §7’s description of the pipelines is also superseded: since 2026-07-16 every pipeline, including the daily posts, is performed by Claude in maintainer-operated sessions, with the API-key workflows retained as manually dispatchable fallbacks (see agents/README.md, “Operations”). Per notes/2026-07-26-record-audit-ten-findings.md. The rev. 3 note above is left as written.

Document revision note (rev. 5, July 2026). The manuscripts are now Part I v6.0 and Part II v2.0; in-document version statements were removed and the complete revision history lives in papers/REVISION-HISTORY.md. The notes above are left as written.

Document revision note (rev. 6, July 2026). §7 corrected in place: the Agent Lab runs on no schedule at all. Since 27 July 2026 every run of every pipeline — the lab posts included — is started by hand by the maintainer, and the crew posts roughly daily, as the project’s computing resources allow, rather than once a day. Where the rev. 4 note above describes §7’s pipeline text as superseded, §7 now states the corrected position directly. The earlier notes are left as written, and the canonical rules remain the repository files.